

# EclairJS

Node.js/JavaScript for Spark

# EclairJS = Node.js + Apache Spark

Why does this equation make sense?

Node.js is considered best for fast, scalable network applications handling large numbers of simultaneous connections

- Non-blocking event loop

However, “Node.js was never created to solve the compute scaling problem. It was created to solve the I/O scaling problem, which it does really well.” (Tomislav Capan) ...

Apache Spark is considered best for fast, scalable compute-intensive applications handling big data, and seamlessly integrating multiple capabilities: streaming, SQL, ML, and graph.

EclairJS enables Node.js developers to program “as-usual” but directly against Spark, taking advantage of its big-data computing and multiple capabilities

# Streaming Air Travel Demonstration

# Code Example - Main Spark Code

```
var spark = require('.././spark.js');

var sc = new spark.SparkContext("spark://107.16.188.227:7077", "Airline Demo");
var sqlContext = new spark.SQLContext(sc);
var ssc;

// data format: rdu,aa,234,sfo,3
function startStream() { StreamingContext will get started with startup of Node.js application.
  ssc = new spark.StreamingContext(sc, new spark.Duration(2000));
  var dstream = spark.KafkaUtils
    .createStream(ssc, "169.54.140.107:2181", "floyd", "airline") Create context for streaming data coming from
    .window(new spark.Duration(1000 * 60 * 15)) Softlayer/Kafka pump with a 15 minute window to
    .flatMap(function(chunk) { check for new data.
      return chunk[1].split('\n');
    })
    .map(function(line) { Read each line from the stream and map it to a
      var lineArr = line.split(","); JSON object of the data we are interested in.
      var str = JSON.stringify({
        "origin": lineArr[16],
        "carrier": lineArr[8],
        "flight_num": lineArr[9],
        "destination": lineArr[17],
        "take_off_delay_mins": parseInt(lineArr[15])
      })
    })

    return str;
  });

  dstream.foreachRDD(function(rdd) { Create a DataFrame for the RDDs coming from the
    // we have a java rdd so wrap it here streamed data that we can perform SQL queries on later.
    var jsRDD = new RDD(rdd);
    if(!jsRDD.isEmpty()) {
      var df = sqlContext.read().json(jsRDD)
      df.registerTempTable("airlinedata")
    }
  }).then(function() {
    ssc.start().catch(function(err) {
      console.log("error starting streaming context");
      console.log(err);
    })
  }).catch(function(err) {
    console.log("error sending print command");
    console.log(err);
  })
}

function getTodayFlights() {
  var file = 'file:' + __dirname + '/public/data/2008bd.json'; On startup of Node.js application we also open
  var dfAllFlights = sqlContext.read().json(file); and read static JSON file to get just flight data for
  var today = new Date(); current day/month. Put into DataFrame tempTable
  var month = today.getMonth()+1; for SQL queries later on.
  var day = today.getDate(); // 1 indexed e.g. 1-31

  var dfFlightsForToday =
    dfAllFlights.filter("month='"+month+"' AND day='"+day+"'");
  dfFlightsForToday.count().then(function(count){
    dfFlightsForToday.registerTempTable('flightstoday');
  });
}
```

# Code Example - Main Node.js Code

```
var express = require('express'); Create a Node.js Express application.
var app = express();

app.use(express.static('public'));
app.use(express.static(__dirname + '/public'));

var airlineDemo = require('./airline.js'); Get an instance for the class that wraps the main Spark code.

app.get('/getFlights', function (request, response) { Handle GET requests for streaming flight data; query the DataFrame where we're dumping flights that have departed within last 15 minutes.
  var airportCode = request.query.airport;
  // data format: rdu,aa,234,sfo,3
  try {
    var df = airlineDemo.query("SELECT * FROM airlinedata WHERE origin='"+airportCode+"'");
  } catch (e) {
    console.log("e", e)
  }

  df.toJSON().toArray().then(function(result) {
    response.json(result);
  }).catch(function(e) {
    console.log(e)
  })
});


app.get('/getCarriers', function (request, response) { Handle GET requests to get carriers for a certain airport; queries the DataFrame with static data for current day's flights.
  var airportCode = request.query.airport;
  try {
    var carriers = airlineDemo.query("SELECT DISTINCT carrier FROM flightstoday WHERE origin='"+airportCode+"'");
    carriers.cache().toJSON().toArray().then(function(result){
      response.json(result);
    });
  } catch (e) {
    console.log("e", e)
  }
});

app.get('/getSchedule', function (request, response) { Handle GET requests to get flight schedule for a given airline carrier at a given airport; queries the DataFrame with static data for current day's flights.
  var airportCode = request.query.airport;
  var carrier = request.query.carrier;
  try {
    var flightsToday = airlineDemo.query("SELECT flight_num,destination FROM flightstoday WHERE origin='"+airportCode + "' AND carrier='"+carrier + "'");
    flightsToday.cache().toJSON().toArray().then(function(result){
      response.json(result);
    });
  } catch (e) {
    console.log("e", e)
  }
});

var server = app.listen(3000, 'localhost', function () { Startup the Node.js application.
  console.log('listening on *:3000');
});


// start the demo
airlineDemo.start(); Tell the class that wraps the Spark code to start the stream and start reading the static data.
```

# EclairJS Node & Nashorn

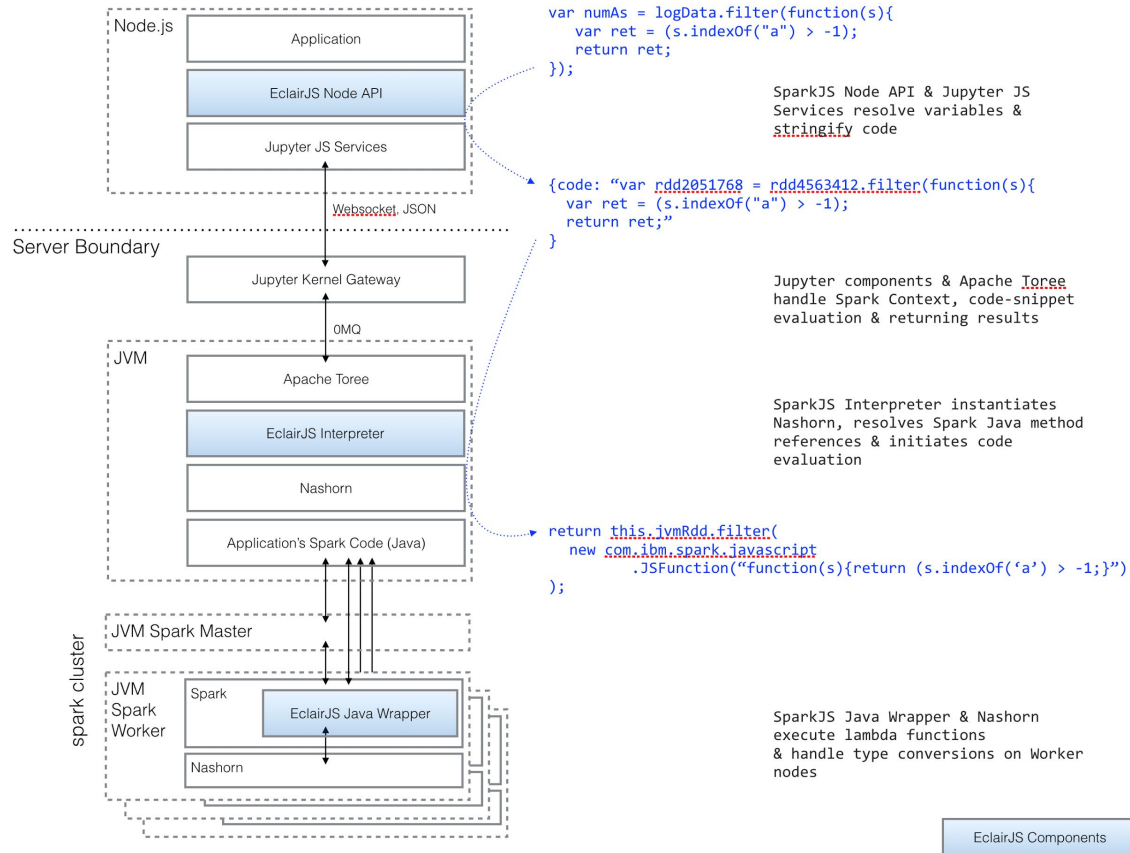


# EclairJS Node

EclairJS Node provides Node.js applications with a Spark API through an npm installable client so that Node.js applications can run remotely from the Spark engine.



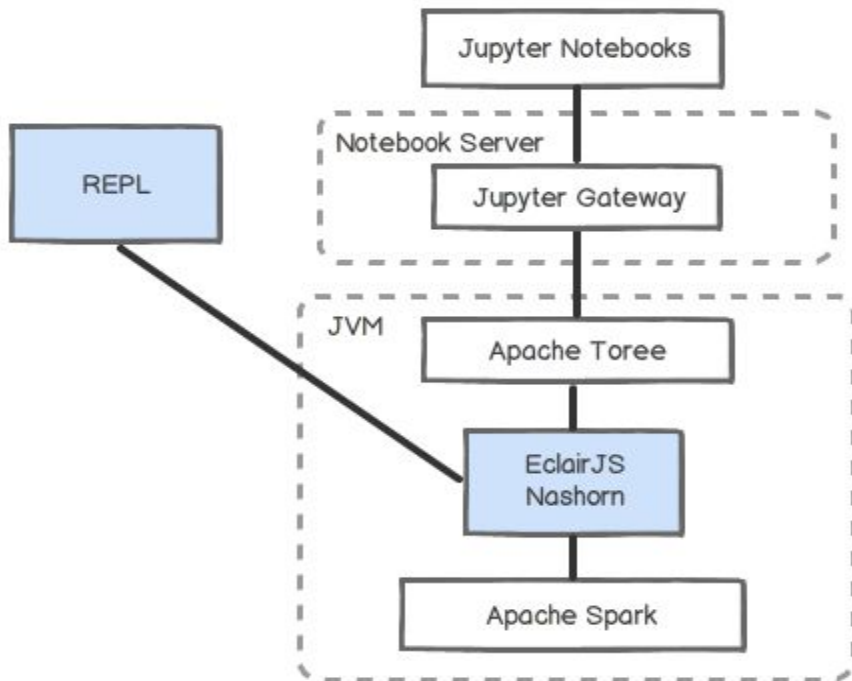
# In-Depth Operation





# EclairJS Nashorn

EclairJS Nashorn implements the support for JavaScript in Spark, and provides a framework that supports various applications including a REPL and Jupyter Notebooks (and EclairJS Node).



```
In [1]: var sparkContext = new SparkContext("local[*]", "myapp");
var sqlContext = new SQLContext(sparkContext);
var peopleDataFrame = sqlContext.read().json("/Users/billreed/eclairjs_dev/eclairjs-nashorn/examples/data/people.json");
var peopleOver20 = peopleDataFrame.filter("age > 20");
eval('"number of people over 20 is " + peopleOver20.count()');
```

```
Out[1]: number of people over 20 is 3
```

```
In [3]: var femalesOver20 = peopleOver20.filter("male = false");
eval('"number of females over 20 is " + femalesOver20.count()');
```

```
Out[3]: number of females over 20 is 1
```

In [ ]:

# Development

Our application target is the cloud.

Our development cycle is:

- Develop Node.js applications locally with local Spark cluster.
- Optional: Test applications against Spark Cluster (with EclairJS Nashorn etc) in SoftLayer
- Push applications to Bluemix cloud where Node.js is available as a service. Provide SoftLayer Spark cluster reference.

# Roadmap

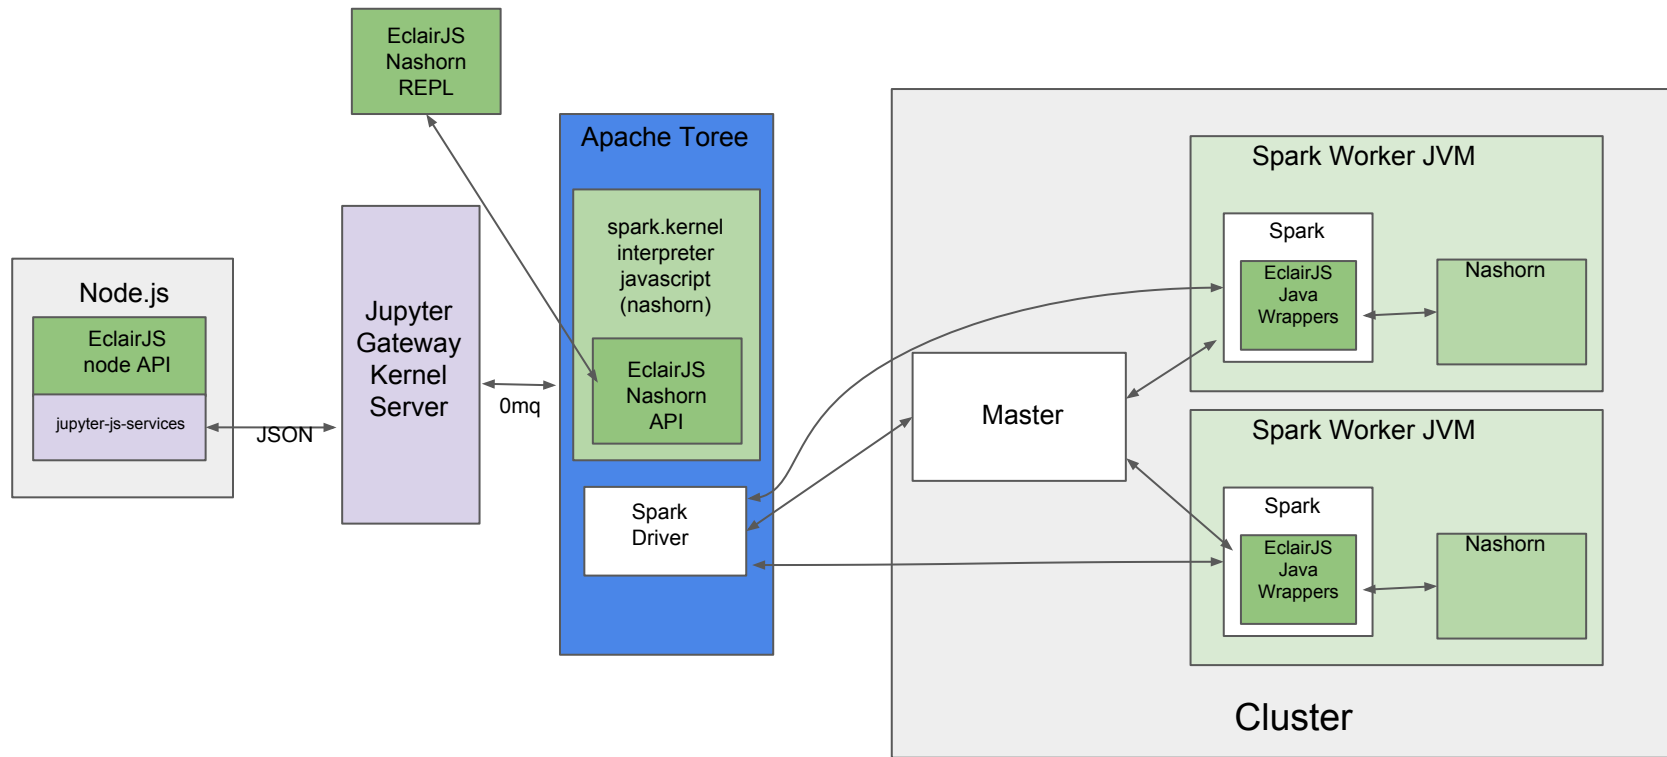
- We are in the process of developing/implementing the entire Spark API.
  - Both **EclairJS** projects are managed in parallel because to develop the *Node* need to develop similar capabilities in the *Nashorn* API, thus the roadmap below components of EclairJS.
- Milestone 1 API support for:
  - Core
  - Streaming
  - SQL/Dataframe
  - Spark 1.6
- Milestone 2 API support for:
  - MLLib
  - GraphX

# Try It Out & Get Involved!

- JavaScript-enabled Jupyter notebooks running in the IBM Bluemix Cloud.
  - Build components from source and setup your own local or clustered environment.
    - Docker Container provides Notebook setup
  - All details are documented on GitHub
  - Projects are open source on GitHub, Apache v2 license
- 
- The projects are looking for contributors and contributions!
  - Google Group - <https://groups.google.com/forum/#!forum/eclairjs>
  - GitHub Project - <https://github.com/EclairJS/eclairjs-node/>, <https://github.com/EclairJS/eclairjs-nashorn>


Back up slides

# EclairJS Node & Nashorn Detail



# Apache Spark<sup>™</sup>

- Fast and general-purpose engine for large-scale data processing
  - Like Hadoop, but using in-memory processing
  - Runs with any Hadoop data source
- Includes libraries for SQL, streaming, machine learning and graph data
  - Can combine multiple libraries within a single application
- Applications can be written in Scala, Java, Python and R
- Scale achieved through parallel distributed-processing: master and workers
- Rapidly being adopted by data engineers for its advanced analytics capabilities, performance, ease of programming
  - Most active Apache project
  - Apache 2 license
  - [spark.apache.org](http://spark.apache.org)





# Jupyter

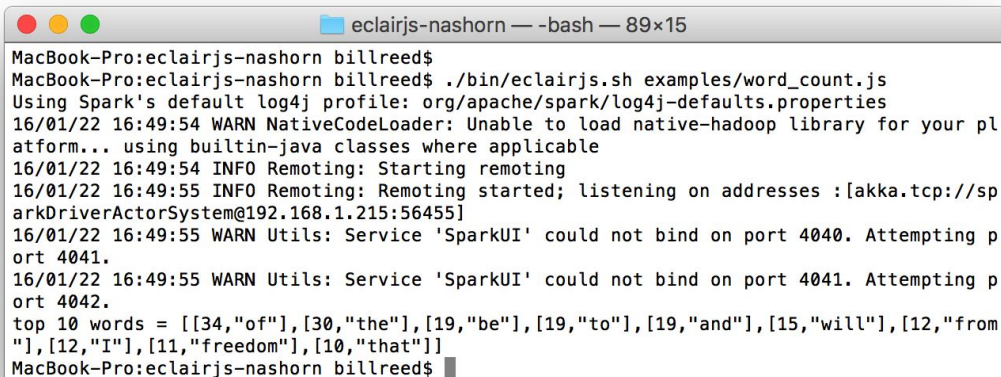
- **Jupyter Notebook** is a web application that allows you to create and share documents that contain live code, equations, visualizations and explanatory text. Uses include: data cleaning and transformation, numerical simulation, statistical modeling, machine learning and much more.
- **Jupyter Kernel Gateway** is a Jupyter application that implements different APIs and protocols for accessing Jupyter kernels such as:
  - Accessing HTTP and Websocket resources of the `/api/kernels` using `jupyter/n`  
`jupyter/jupyter_client` and `jupyter/jupyter_core`
  - Accessing notebook cells via HTTP endpoints

# EclairJS REPL Example

```
eclairjs-nashorn — java -cp /usr/local/spark-1.6.0-bin-hadoop2.6/conf/:/usr/local/spark-1.6.0-bin-hadoop2.6/lib/spark-assembly-1.6.0-hadoop2.6.0.jar:/usr/l...
MacBook-Pro:eclairjs-nashorn billreed$ ./bin/eclairjs.sh
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
16/01/22 16:24:28 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
16/01/22 16:24:28 INFO Remoting: Starting remoting
16/01/22 16:24:28 INFO Remoting: Remoting started; listening on addresses :[akka.tcp://sparkDriverActorSystem@192.168.1.215:56340]
16/01/22 16:24:29 WARN Utils: Service 'SparkUI' could not bind on port 4040. Attempting port 4041.
16/01/22 16:24:29 WARN Utils: Service 'SparkUI' could not bind on port 4041. Attempting port 4042.
Welcome to eclairJS-nashorn, Type in expressions to have them evaluated.
SQL context available as sc..
eclairjs>var sqlContext = new SQLContext(sc);
null
eclairjs>var peopleDataFrame = sqlContext.read().json("/Users/billreed/eclairjs_dev/eclairjs-nashorn/examples/data/people.json");
null
eclairjs>var peopleOver20= peopleDataFrame.filter("age > 20");
null
eclairjs>peopleOver20.show();
+-----+
|age|male|name|
+-----+
|29|true|Michael|
|40|true|Andy|
|33|false|Sue|
+-----+

null
eclairjs>
```

# EclairJS JavaScript Driver Example



```
eclairjs-nashorn — -bash — 89x15
MacBook-Pro:eclairjs-nashorn billreed$
MacBook-Pro:eclairjs-nashorn billreed$ ./bin/eclairjs.sh examples/word_count.js
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
16/01/22 16:49:54 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
16/01/22 16:49:54 INFO Remoting: Starting remoting
16/01/22 16:49:55 INFO Remoting: Remoting started; listening on addresses :[akka.tcp://sparkDriverActorSystem@192.168.1.215:56455]
16/01/22 16:49:55 WARN Utils: Service 'SparkUI' could not bind on port 4040. Attempting port 4041.
16/01/22 16:49:55 WARN Utils: Service 'SparkUI' could not bind on port 4041. Attempting port 4042.
top 10 words = [[34,"of"],[30,"the"],[19,"be"],[19,"to"],[19,"and"],[15,"will"],[12,"from"],[12,"I"],[11,"freedom"],[10,"that"]]
MacBook-Pro:eclairjs-nashorn billreed$
```

```
var file = "src/test/resources/dream.txt";
var conf = new SparkConf()
    .setAppName("JavaScript word count")
    .setMaster("local[*]");
var sparkContext = new SparkContext(conf);
var rdd = sparkContext.textFile(file).cache();
var rdd2 = rdd.flatMap(function(sentence) {
    return sentence.split(" ");
});
var rdd3 = rdd2.filter(function(word) {
    return word.trim().length > 0;
});
var rdd4 = rdd3.mapToPair(function(word) {
    return [word, 1];
});
var rdd5 = rdd4.reduceByKey(function(a, b) {
    return a + b;
});
var rdd6 = rdd5.mapToPair(function(tuple) {
    return [tuple[1]+0.0, tuple[0]];
});
var rdd7 = rdd6.sortByKey(false);
print("top 10 words = " + JSON.stringify(rdd7.take(10)));
sparkContext.stop();
```